

# Day 3: Fundamentals of Deep Learning

Naeemullah Khan

[naeemullah.khan@kaust.edu.sa](mailto:naeemullah.khan@kaust.edu.sa)



جامعة الملك عبد الله  
للعلوم والتكنولوجيا  
King Abdullah University of  
Science and Technology

KAUST Academy  
King Abdullah University of Science and Technology

December 28, 2025

*By the end of this session, you will be able to:*

1. Define Deep Learning and explain why it scales better than traditional Machine Learning.
2. Dissect a Neural Network into its core components: Neurons, Layers, and Weights.
3. Trace the Learning Process, explaining how data flows forward and gradients flow backward (Backpropagation).
4. Compare Optimizers (SGD, Momentum, Adam) and understand how they work.

# Table of Contents

1. Introduction to Deep Learning
2. Anatomy of a Neural Network
3. Training a Neural Network
4. A Deep Dive into Optimizers
5. What We Choose When Building a Neural Network
6. Types of Neural Network Architectures
7. References

# What is Deep Learning?

Deep Learning is simply **Machine Learning**...

...but powered by **Artificial Neural Networks** with many layers.

- ▶ "Neural": Inspired by the human brain (neurons).
- ▶ "Deep": Refers to the number of layers (depth).

# What is Deep Learning?

Deep Learning is simply **Machine Learning**...

...but powered by **Artificial Neural Networks** with many layers.


- ▶ "Neural": Inspired by the human brain (neurons).
- ▶ "Deep": Refers to the number of layers (depth).

*But why did Deep Learning suddenly become so powerful?*

# Why Now?

*Deep Learning isn't new. The ideas date back decades. But three things changed:*


1. More **data**  
(internet, sensors, images)
2. Faster **compute**  
(GPUs, cloud computing)
3. Better **algorithms**  
(training techniques,  
architectures)



DL keeps improving as data grows.  
ML hits a ceiling.

*Deep Learning isn't new. The ideas date back decades. But three things changed:*

1. More **data**  
(internet, sensors, images)
2. Faster **compute**  
(GPUs, cloud computing)
3. Better **algorithms**  
(training techniques,  
architectures)



DL keeps improving as data grows.  
ML hits a ceiling.

*But what makes Deep Learning so scalable?*

# The Secret: Simple Building Blocks

## Traditional ML

Models are **fixed equations** designed by experts with specific assumptions about the data.

*Example:* Linear Regression  $y = w_1x + w_0$


**Benefit:** Fast and interpretable (when assumptions hold).

**Problem:** Hard to adapt when data nature is different or unknown.

## Deep Learning

Models are built from **simple, stackable blocks** (like LEGO).

*Example:* Stacking layers



**Benefit:** Very flexible. Can be adapted to any data type or task.

# The Secret: Simple Building Blocks

## Traditional ML

Models are **fixed equations** designed by experts with specific assumptions about the data.

*Example:* Linear Regression  $y = w_1x + w_0$


**Benefit:** Fast and interpretable (when assumptions hold).

**Problem:** Hard to adapt when data nature is different or unknown.

## Deep Learning

Models are built from **simple, stackable blocks** (like LEGO).



*Example:* Stacking layers




**Benefit:** Very flexible. Can be adapted to any data type or task.

*So what is this magic building block?*

The magic building block is called a Neuron.



The magic building block is called a Neuron.



*But how powerful can these simple blocks really be?*


# Surprisingly Powerful: Universal Approximation

Here's the remarkable part: if you stack enough neurons together, they can learn to fit **any pattern**, no matter how simple/complex it is.

## Universal Approximation Theorem

A Neural Network with enough neurons can approximate **any continuous function** to arbitrary accuracy.

This is why neural networks are so flexible and powerful!



*"Give me enough neurons,  
and I can fit any function."*


# Surprisingly Powerful: Universal Approximation

Here's the remarkable part: if you stack enough neurons together, they can learn to fit **any pattern**, no matter how simple/complex it is.

## Universal Approximation Theorem

A Neural Network with enough neurons can approximate **any continuous function** to arbitrary accuracy.

This is why neural networks are so flexible and powerful!



*"Give me enough neurons,  
and I can fit any function."*

*So let's look inside: what does a neural network actually look like?*

# Inside a Neural Network


A Neural Network is organized into **layers** of neurons, where each layer transforms the data step by step.

## The Three Parts:

- ▶ **Input Layer:** Receives your raw data ( $x$ ).
- ▶ **Hidden Layers:** Extract patterns and features.
- ▶ **Output Layer:** Produces the final prediction ( $\hat{y}$ ).


The "deep" in Deep Learning comes from having many hidden layers.

Input      Hidden      Output



*Each circle is a neuron,  
each line is a connection.*

# Example: Predict House Price




A Neural Network is organized into **layers** of neurons, where each layer transforms the data step by step.

## The Three Parts:

- ▶ **Input Layer:** Receives your raw data ( $x$ ).
- ▶ **Hidden Layers:** Extract patterns and features.
- ▶ **Output Layer:** Produces the final prediction ( $\hat{y}$ ).

The "deep" in Deep Learning comes from having many hidden layers.



*Each circle is a neuron,  
each line is a connection.*

*Now let's zoom in and see what a single neuron actually does...*

# What Does a Neuron Do?

A neuron performs two simple operations:

## 1. Linear Combination:

$$z = w^\top x + b$$


(Just like linear regression!)

## 2. Non-linear Activation:

$$a = \sigma(z)$$

(This is the new part)

The output  $a$  becomes the input for the next layer neurons!



# What Does a Neuron Do?

A neuron performs two simple operations:

## 1. Linear Combination:

$$z = w^\top x + b$$


(Just like linear regression!)

## 2. Non-linear Activation:

$$a = \sigma(z)$$

(This is the new part)

The output  $a$  becomes the input for the next layer neurons!



*But why do we need that activation function  $\sigma$ ? Why don't we just use the linear part ( $wx + b$ )?*

# Why We Need Activation Functions

**Question:** Why don't we just use the linear part ( $wx + b$ )?

## The Problem: Linearity

Without activations, stacking linear layers gives us... just another linear function!

$$\text{Linear} + \text{Linear} = \text{Linear}$$

No matter how many layers you stack, without activations, the whole network collapses into a single Linear Regression model!

# Why We Need Activation Functions

**Question:** Why don't we just use the linear part ( $wx + b$ )?

## The Problem: Linearity

Without activations, stacking linear layers gives us... just another linear function!

$$\text{Linear} + \text{Linear} = \text{Linear}$$


No matter how many layers you stack, without activations, the whole network collapses into a single Linear Regression model!

**Solution: The "Bend":** Activations ( $\sigma$ ) introduce non-linearity. Think of it like bending a ruler. This allows us to fit curved, complex patterns.

# Common Activation Functions

**ReLU**


$$\sigma(z) = \max(0, z)$$



**Most popular!** Simple and works well in deep networks.

**Sigmoid**


$$\sigma(z) = \frac{1}{1 + e^{-z}}$$



Output in  $(0, 1)$ . You know this from logistic regression!

**Tanh**


$$\sigma(z) = \tanh(z)$$



Output in  $(-1, 1)$ . Similar to sigmoid but centered at zero.

# From Neurons to Layers

We don't just use one neuron. We stack them in parallel to form a **Layer**.



## How do we decide the number of neurons?

- ▶ **The Input Layer:** If our data has  $k$  features (e.g., Age, Salary, Debt), the input layer must have  $k$  neurons.
- ▶ **The Output Layer:** If we are doing multiclass classification (e.g., Cat, Dog, Bird), the output layer must have  $k$  neurons (one score per class).
- ▶ **Hidden Layers:** This is a design choice (Hyperparameter). More neurons = captures more complex patterns.

# Network Shapes for Different Tasks

Task	Input	Output of NN	Typical loss
Single output regression	$x \in \mathbb{R}^k$	$\hat{y} \in \mathbb{R}$	Mean squared error
Multi output regression	$x \in \mathbb{R}^k$	$\hat{y} \in \mathbb{R}^m$	MSE over all outputs
Binary classification	$x \in \mathbb{R}^k$	$\hat{p} \in (0, 1)$	Binary cross entropy
Multiclass classification	$x \in \mathbb{R}^k$	$\hat{p} \in \mathbb{R}^C$ softmax	Cross entropy

Since a layer is just a collection of neurons running in parallel, we can mathematically represent it using **Linear Algebra**.

► **Single Neuron:**

$$z = \mathbf{w}^\top \mathbf{x} + b$$

(Weights are a vector  $\mathbf{w}$ )

► **Full Layer:**

$$\mathbf{Z} = \mathbf{W}\mathbf{x} + \mathbf{b}$$

(Weights are a **Matrix  $\mathbf{W}$** )


*Every row in matrix  $\mathbf{W}$  represents one specific neuron's weights.*

# From Layer to Layers (Going Deep)

**In theory:** One layer is enough to fit anything.


**In practice:** Deeper networks are much more powerful and efficient.

- ▶ They learn **hierarchical features**:  
(Pixel → Edge → Shape → Face).



# Putting it Together: The Neural Network

A Neural Network is just a **recursive chain** of these layers. The output of one layer becomes the input of the next.




**Mathematically (The Recursive Formula):**

1. **Input:**  $h^{(0)} = x$
2. **Hidden 1:**  $h^{(1)} = \sigma(W_1 h^{(0)} + b_1)$
3. **Hidden 2:**  $h^{(2)} = \sigma(W_2 h^{(1)} + b_2)$
4. **Output:**  $\hat{y} = \text{Softmax}(W_3 h^{(2)} + b_3)$

*Data flows from left to right, getting transformed at every step.*

# Neural Networks as Part of a Learning System



# Training a Neural Network

To train this model  
we repeat two phases:


## Forward Pass:

- ▶ Data → Model  
→ Predictions
- ▶ Compute Loss

## Backward Pass:

- ▶ Calculate  
gradients
- ▶ Update  
parameters

*Let's discuss them in  
detail.*



# 1. Forward Pass

**Goal:** Compute the network output  $\hat{y}$  from input  $x$ , then measure the error.

**What happens (layer by layer):**


1. Start with input:  $a^{(0)} = x$ .
2. For each layer  $\ell$ :

$$z^{(\ell)} = W^{(\ell)} a^{(\ell-1)} + b^{(\ell)}$$

$$a^{(\ell)} = \sigma(z^{(\ell)})$$

3. Final layer gives the prediction:  
 $\hat{y} = a^{(L)}$ .
4. **Loss:**  $J(\theta)$  compares  $\hat{y}$  with the true target  $y$ .

*Forward pass = just prediction (no learning yet).*




## 2. Backward Pass (Backpropagation)

**Goal:** Given feedback from the loss, calculate gradients for all weights, then update them via Gradient Descent.

By applying the **Chain Rule**, we multiply these partial derivatives to derive the gradient for every parameter.

We move from the end (the loss) **backwards** towards the beginning, passing through each weight to get its gradient. This is called **Backpropagation**.



# Chain Rule Example: Compute $\frac{\partial J}{\partial w}$

Assume one neuron:

$$z = wx + b, \quad a = \sigma(z), \quad J = \frac{1}{2}(a - y)^2$$

## Chain Rule (backprop)

$$\frac{\partial J}{\partial w} = \frac{\partial J}{\partial a} \cdot \frac{\partial a}{\partial z} \cdot \frac{\partial z}{\partial w}$$

Now compute each piece:

$$\frac{\partial J}{\partial a} = (a - y) \quad \frac{\partial a}{\partial z} = \sigma(z)(1 - \sigma(z)) = a(1 - a) \quad \frac{\partial z}{\partial w} = x$$

Final result

$$\frac{\partial J}{\partial w} = (a - y) a(1 - a) x$$

**What do we do with  $\frac{\partial J}{\partial w}$  ?** We use Gradient Descent to update the weights.

## Parameter Update Rule (via Gradient Descent)

$$w \leftarrow w - \alpha \frac{\partial J}{\partial w} \quad b \leftarrow b - \alpha \frac{\partial J}{\partial b}$$

- ▶  $\alpha$  : learning rate (step size).
- ▶ We compute  $\frac{\partial J}{\partial(\cdot)}$  for **every** parameter using backprop.

**What do we do with  $\frac{\partial J}{\partial w}$  ?** We use Gradient Descent to update the weights.

## Parameter Update Rule (via Gradient Descent)

$$w \leftarrow w - \alpha \frac{\partial J}{\partial w} \quad b \leftarrow b - \alpha \frac{\partial J}{\partial b}$$


- ▶  $\alpha$  : learning rate (step size).
- ▶ We compute  $\frac{\partial J}{\partial(\cdot)}$  for **every** parameter using backprop.

**Q:** Do we do this update using **all data at once?**

# Training Step = One Update

We do not train using the full dataset at once.

Instead, we take a **small batch** (mini-dataset) and do one update:




- ▶ This single update is called a **training step** (or **iteration**).
- ▶ **Batch size** = how many samples are used in one step.

*Next: repeat this step for all batches ⇒ one full epoch.*

# Epoch = One Full Pass Over the Dataset

Dataset → split into **batches** → each batch makes **one step**.



- ▶ **Steps per epoch**  $\approx \left\lceil \frac{\# \text{samples}}{\text{batch size}} \right\rceil$ .
- ▶ Training usually runs for **many epochs**.

*Next: before the next epoch, we usually shuffle the dataset.*

## Why Shuffle Between Epochs?


**At the start of each epoch, we shuffle the dataset to randomize batch composition.**

## Without shuffling:

- ▶ Batches see same patterns repeatedly
  - ▶ Similar samples grouped together
  - ▶ Biased gradient estimates

## With shuffling:

- ▶ Each batch is a better mix
  - ▶ More stable learning
  - ▶ Better generalization



# The Three Pillars of Learning

## 1. The Model ( $f_\theta$ )

Design Choice:  
Neural Networks  
(Architecture)

## 2. Loss Function ( $J$ )

Design Choice:  
MSE / Cross-Entropy

## 3. Optimizer

Design Choice:  
SGD, Adam, RMSProp


Depends on:  
Stability & Convergence Speed

We defined the Models (1) and the Losses (2).  
Now, Let's discuss optimizers.

# Gradient Descent: The Core Idea

**Goal:** Find parameters  $\theta$  that minimize the loss  $J(\theta)$ .

**Intuition:** Follow the negative gradient (steepest descent direction).



The gradient  $\nabla_{\theta} J(\theta)$  tells us which direction increases the loss most.  
We go the **opposite direction** to decrease it.

# Gradient Descent: Update Rule

**Basic update equation:**

$$\theta_{\text{new}} = \theta_{\text{old}} - \alpha \nabla_{\theta} J(\theta)$$

- ▶  $\alpha$  is the **learning rate** (step size)
- ▶  $\nabla_{\theta} J(\theta)$  is the gradient of loss w.r.t. parameters
- ▶ Repeat until convergence (or for fixed number of steps)

# Gradient Descent: Update Rule

**Basic update equation:**


$$\theta_{\text{new}} = \theta_{\text{old}} - \alpha \nabla_{\theta} J(\theta)$$

- ▶  $\alpha$  is the **learning rate** (step size)
- ▶  $\nabla_{\theta} J(\theta)$  is the gradient of loss w.r.t. parameters
- ▶ Repeat until convergence (or for fixed number of steps)

**Three variants based on how much data we use:**

1. **Batch Gradient Descent:** Use entire dataset
2. **Stochastic Gradient Descent (SGD):** Use one sample
3. **Mini-Batch Gradient Descent:** Use small batch of samples

How much data do we use per update (one step)?



## Batch GD

- ▶ Stable updates
- ▶ Expensive per step
- ▶ Needs lots of RAM

## Mini-Batch (Default)

- ▶ Best of both worlds
- ▶ Stable & Fast
- ▶ GPU Efficient


## Stochastic GD

- ▶ Very noisy updates
- ▶ Can escape local minima
- ▶ Slow (can't parallelize)


*In practice: we almost always use Mini-Batch (whenever GD mentioned next, it refers to the mini-batch variant).*

# Convergence Behavior

How do they navigate the Loss Surface?



**The Ravine Problem:** Gradient Descent makes slow progress because it "bounces" off the steep walls.



## The Solution

We add **Momentum**, which acts like a heavy ball rolling downhill, building velocity and ignoring small bumps.

It maintains a **running average** of past gradients ( $v_t$ ), allowing oscillations in opposing directions to cancel out while consistent directions accumulate and accelerate.

Keep a moving average of past gradients:


$$v_t = \underbrace{\beta v_{t-1}}_{\text{past gradients}} + \underbrace{(1 - \beta) \nabla_{\theta} J(\theta)}_{\text{new gradient}}$$

$$\theta_{t+1} = \theta_t - \alpha v_t$$

- ▶  $v_t$  is the **velocity** (momentum term)
- ▶  $\beta$  is the momentum coefficient (typically 0.9)
- ▶  $\alpha$  is the learning rate

**Problem:** Same learning rate for all parameters may not be optimal.

**Idea:** Give each parameter its own adaptive learning rate based on its gradient history.



Parameters with large gradients get smaller steps (learning rates).  
This prevents overshooting.

**Accumulate squared gradients for each parameter:**

$$G_t = \underbrace{G_{t-1}}_{\text{past squared gradients}} + \underbrace{g_t^2}_{\text{new squared gradient}}$$


**Update with parameter-wise learning rate:**

$$\theta_{t+1} = \theta_t - \underbrace{\frac{\eta}{\sqrt{G_t} + \epsilon}}_{\text{adaptive learning rate}} \underbrace{g_t}_{\text{gradient}}$$

- ▶  $G_t$  accumulates all past squared gradients  $\rightarrow$  larger  $G_t$  means smaller effective learning rate
- ▶  $\eta$  is the base learning rate
- ▶  $\epsilon$  is a small constant for numerical stability ( $\sim 10^{-8}$ )
- ▶ Each parameter gets its own effective learning rate

**Problem with AdaGrad:** Accumulates *all* past gradients → Learning rate can shrink too much over time (never forgets past gradients).

**RMSProp solution:** Use **exponential moving average** instead (forget old gradients!).



**Effect:** Learning rate adapts but doesn't vanish → training continues effectively!

# RMSProp: Update Rule

**Exponential moving average of squared gradients:**

$$E[g^2]_t = \underbrace{\alpha E[g^2]_{t-1}}_{\text{decayed past gradients}} + \underbrace{(1 - \alpha)g_t^2}_{\text{new squared gradient}}$$


**Update rule:**

$$\theta_{t+1} = \theta_t - \frac{\eta}{\sqrt{E[g^2]_t + \epsilon}} \underbrace{g_t}_{\text{gradient}}$$

adaptive learning rate

- ▶  $\alpha$  is the decay rate (typically 0.9) (old gradients gradually fade away)
- ▶ Recent gradients have more influence than old ones

**Idea:** Combine **momentum** + **adaptive learning rates!**



**Bias Correction:** At the start of training, moving averages are initialized to zero, causing them to be biased toward zero. Adam corrects this bias to get accurate estimates in early steps.

**Result:** Fast, stable, and works well across many problems (the most popular optimizer in deep learning!)

# Adam: Update Rule

Maintain two moving averages:

$$m_t = \underbrace{\beta_1 m_{t-1}}_{\text{past momentum}} + \underbrace{(1 - \beta_1) g_t}_{\text{new gradient}} \quad (\text{momentum})$$

$$v_t = \underbrace{\beta_2 v_{t-1}}_{\text{past variance}} + \underbrace{(1 - \beta_2) g_t^2}_{\text{new squared gradient}} \quad (\text{adaptive LR})$$

**Bias correction** (fixes initial zero bias):

$$\hat{m}_t = \frac{m_t}{1 - \beta_1^t}, \quad \hat{v}_t = \frac{v_t}{1 - \beta_2^t}$$

As  $t \rightarrow \infty$ ,  $\beta^t \rightarrow 0$ , so correction factor  $\rightarrow 1$  (no correction needed later)


**Update:**

$$\theta_{t+1} = \theta_t - \underbrace{\frac{\eta}{\sqrt{\hat{v}_t} + \epsilon}}_{\text{adaptive step size}} \underbrace{\hat{m}_t}_{\text{momentum direction}}$$

► Typical:  $\beta_1 = 0.9$ ,  $\beta_2 = 0.999$

**Problem with Adam:** Weight decay (L2 regularization) interacts poorly with adaptive learning rates.

**Why?** In Adam, weight decay is added to the gradient before the adaptive scaling. This means parameters with small gradients get disproportionately large weight decay, breaking the intended regularization.



**Result:** Better generalization, especially for large models (current best practice for training modern neural networks!)


Same as Adam, plus decoupled weight decay:

$$\theta_{t+1} = \theta_t - \underbrace{\frac{\eta}{\sqrt{\hat{v}_t + \epsilon}} \hat{m}_t}_{\text{Adam update}} - \underbrace{\eta \lambda \theta_t}_{\text{weight decay (decoupled)}}$$

- ▶ First term: standard Adam update (gradient-based optimization)
- ▶ Second term: direct weight decay applied to parameters, independent of gradients
- ▶  $\lambda$  is the weight decay coefficient (typically 0.01)

**Why better?** Cleaner separation of optimization and regularization.

## Evolution of Optimizers:



**Rule of thumb:** Start with **AdamW** (it's the most robust default choice!)

## Key Design Decisions:

### 1. Model Architecture

#### Structure:

- ▶ Number of layers
- ▶ Neurons per layer
- ▶ Network type  
(MLP/CNN/RNN)

#### Components:

- ▶ Activation functions  
(ReLU, sigmoid, tanh)
- ▶ Output layer shape & activation

### 2. Training Strategy

#### Learning:


- ▶ Loss function (MSE, cross-entropy)
- ▶ Optimizer (SGD, Adam, AdamW)

#### Hyperparameters:


- ▶ Learning rate
- ▶ Batch size
- ▶ Number of epochs

# Types of Neural Network Architectures


- ▶ **MLP** fully connected network for tabular data
- ▶ **CNN** convolutional neural network for images
- ▶ **RNN and variants** for sequences and time series



Standard NN



Convolutional NN



Recurrent NN

- ▶ Aurélien Géron, *Hands-On Machine Learning with Scikit-Learn, Keras, and TensorFlow*
- ▶ Andrew Ng, *Deep Learning Specialization* (Coursera/DeepLearning.AI)

*Slides contributed by Mohamed Eltayeb*